

Practices in visual computing 1

Lab2: Numpy & OpenCV

Simon Fraser University

Fall 2025

Numerical Computing with NumPy

Foundation for scientific Python stack

Powers Machine Learning, Computer Vision, and data analysis workflows



Why NumPy?

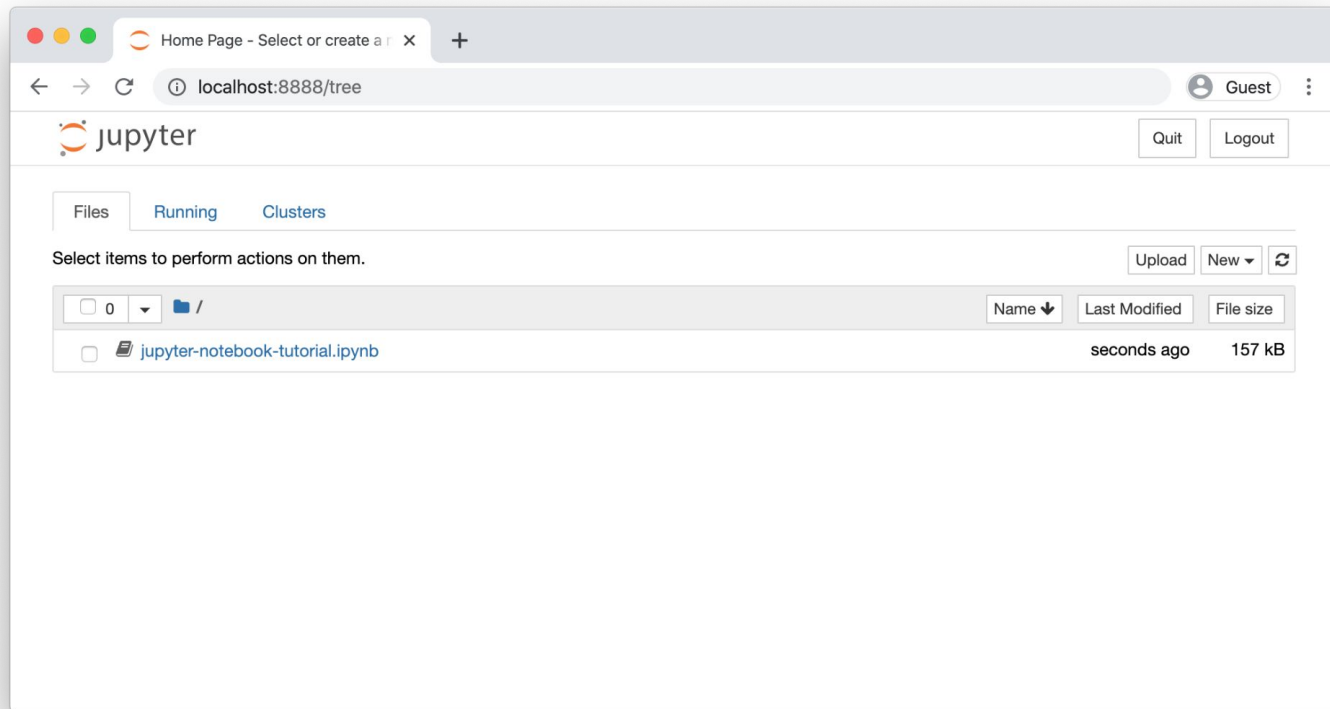
Speed: implemented in C under the hood

Memory efficiency: contiguous storage

Vectorized operations (no explicit loops)

Integration: works seamlessly with OpenCV, SciPy, scikit-learn, etc.

Jupyter and Colab Notebooks



NumPy Arrays vs. Python Lists

Python Lists:

- Store mixed types
- Extra overhead (each element = Python object)
- Scattered in memory → slower math

NumPy Arrays:

- Homogeneous (one data type)
- Contiguous memory → compact & fast
- Backed by efficient C implementation

NumPy Views

A view is a new array object that shares the same underlying data.

No extra memory is used → slicing, reshaping, transposing are fast.

Changing data in a view also changes it in the original array.

View vs. Copy

2D [4x3]
NumPy ndarray

0	1	2
3	4	5
6	7	8
9	10	11

Flags

C_CONTIGUOUS: True
F_CONTIGUOUS: False
OWNDATA: True

itemsize: 8byte
shape: (4, 3)
strides: (24, 8)
base: None
data: Memory View #1

Transpose(x)

0	3	6	9
1	4	7	10
2	5	8	11

Flags

C_CONTIGUOUS: False
F_CONTIGUOUS: True
OWNDATA: False

itemsize: 8byte
shape: (3, 4)
strides: (8, 24)
base: x
data: Memory View #1

same shape
same data
different strides!

Transpose(x).Copy()

0	3	6	9
1	4	7	10
2	5	8	11

Flags

C_CONTIGUOUS: True
F_CONTIGUOUS: False
OWNDATA: True

itemsize: 8byte
shape: (3, 4)
strides: (32, 8)
base: None
data: Memory View #2

Memory View #1

Memory shared with x

0	1	2	3	4	5	6	7	8	9	10	11
---	---	---	---	---	---	---	---	---	---	----	----

8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte

Memory View #2

Different Memory Layout

0	3	6	9	1	4	7	10	2	3	8	11
---	---	---	---	---	---	---	----	---	---	---	----

8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte 8byte

Some array creation routines


Numerical ranges: `arange`, `linspace`, `logspace`

Homogeneous data: `zeros`, `ones`


Diagonal elements: `diag`, `eye`

Random numbers: `rand`, `randint`



```
>>> np.arange(1, 10, 3)  
array([1, 4, 7])
```



```
>>> np.arange(1, 8, 3)  
array([1, 4, 7])
```




```
>>> np.arange(1, 10.1, 3)  
array([1., 4., 7., 10.])
```





```
np.linspace(2, 10, 5)
```



Equally spaced values

Indexing and slicing in one dimension

1d arrays: indexing and slicing as for lists

- first element has index 0
- negative indices count from the end
- slices: [start:stop:step] without the element indexed by stop
- if values are omitted:
 - start: starting from first element
 - stop: until (and including) the last element
 - step: all elements between start and stop-1

Indexing and slicing in higher dimensions

`a[2, -3]`

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

$a[:3, :5]$

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

YOUR TURN

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

$a[-3:, -3:]$

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

YOUR TURN

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

$a[:, 3]$

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

YOUR TURN

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

`a[1, 3:6]`

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

YOUR TURN

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

`a[1::2, ::3]`

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

YOUR TURN

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Indexing and slicing in higher dimensions

$a[a \% 3 == 0]$

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Fancy indexing – array of integers

$a[(1, 1, 2, 2, 3, 3), (3, 4, 2, 5, 3, 4)]$

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39

Matrix multiplication

$$\begin{pmatrix} 0 & 1 \\ 2 & 3 \end{pmatrix} \begin{pmatrix} 4 & 5 \\ 6 & 7 \end{pmatrix} = \begin{pmatrix} 6 & 7 \\ 26 & 31 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1 \\ 2 & 3 \end{pmatrix} \begin{pmatrix} 4 & 5 \\ 6 & 7 \end{pmatrix} = \begin{pmatrix} 6 & 7 \\ 26 & 31 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1 \\ 2 & 3 \end{pmatrix} \begin{pmatrix} 4 & 5 \\ 6 & 7 \end{pmatrix} = \begin{pmatrix} 6 & 7 \\ 26 & 31 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1 \\ 2 & 3 \end{pmatrix} \begin{pmatrix} 4 & 5 \\ 6 & 7 \end{pmatrix} = \begin{pmatrix} 6 & 7 \\ 26 & 31 \end{pmatrix}$$

```
try np.dot(●, ●)
    ●.dot(●)
    ●@●*)
```

*) Python ≥ 3.5 , NumPy ≥ 1.10

Sorting, searching, and counting in NumPy

Sorting:

sort, argsort, ndarray.sort, sort_complex, partition, argpartition

Searching:


argmax, nanargmax, argmin, nanargmin, argwhere, nonzero, where, searchsorted

Counting:

count_nonzero

K-Means Clustering

Goal: Group data into k clusters by similarity



Algorithm Steps

Initialize:

Choose k cluster centers (random or k-means++).

Assign:

Each point goes to the nearest center.


Update:

Recompute centers as the mean of assigned points.


Repeat:

Alternate assign/update until centers stop moving.


Algorithm Steps




1- Initialize




2- Assign



3- Update



4- Repeat 2 & 3




Weakness of naive initialization

In basic K-Means, the centroids are initialized randomly.

Bad initialization can cause:

- Slow convergence (more iterations).
- Poor clustering (algorithm stuck in bad local minima).



K-Means++ Initialization Strategy

Instead of picking all centroids at random, K-Means++ spreads them out intelligently:

1. Pick the first centroid randomly from the data.
2. For each remaining point, compute its distance to the nearest chosen centroid.
3. Pick the next centroid with probability proportional to the squared distance.

Farther points are more likely to be picked.

4. Repeat until k centroids are chosen.
5. Run normal K-Means as usual.

Introduction to OpenCV

Open Source Computer Vision Library

Focus: image processing, computer vision, machine learning

Widely used in academia & industry



Images in OpenCV

Images are NumPy arrays ($H \times W \times C$)

Color channels: BGR (not RGB)

Loading:

```
# Importing the OpenCV Library  
import cv2  
# Reading the image using imread() function  
image = cv2.imread('image.jpg')  
  
# Extracting the height and width of an image  
h, w = image.shape[:2]  
# Displaying the height and width  
print("Height = {}, Width = {}".format(h, w))
```



Basic Operations

Resize: `cv2.resize()`

Crop: NumPy slicing

Rotate & flip: `cv2.rotate()`, `cv2.flip()`

Convert color: `cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)`

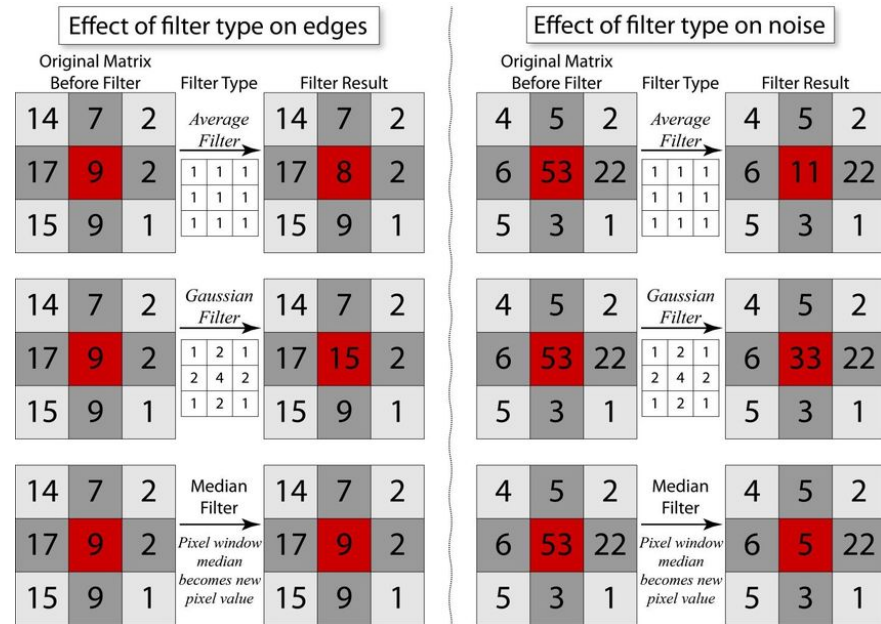


Filtering & Blurring


Smoothing reduces noise

Common filters:


- Average blur → `cv2.blur()`
- Gaussian blur → `cv2.GaussianBlur()`
- Median blur → `cv2.medianBlur()`



Effect of filter type on edges



Effect of filter type on noise



Edge Detection

Edges = sudden changes in intensity (where brightness or color shifts quickly).

Important for:

1. Object boundaries
2. Feature extraction
3. Shape analysis



Sobel Operator

Approximates the image derivative → highlights intensity changes (edges).

Uses 3×3 kernels separately for x and y directions.

X-gradient kernel: negatives on the left, positives on the right → detects vertical edges.

Y-gradient kernel: negatives on the bottom, positives on the top → detects horizontal edges.

-1	0	+1
-2	0	+2
-1	0	+1

G_x

+1	+2	+1
0	0	0
-1	-2	-1

G_y

Sobel Operator

Place a gradient kernel (3×3) over each pixel of the image.

Compute difference in intensity \rightarrow approximates the derivative.

Produces two output images:

- X-Direction (G_x): vertical edges
- Y-Direction (G_y): horizontal edges

Using kernel convolution, sharp changes appear as bright lines (edges).

Sobel Operator

100	100	200	200
100	100	200	200
100	100	200	200
100	100	200	200

-1	0	1
-2	0	2
-1	0	1

$$\begin{array}{r} -100 \\ -200 \\ -100 \\ 200 \\ 400 \\ \hline +200 \\ \hline =400 \end{array}$$

Kernel Convolution: The bigger the value at the end, the more noticeable the edge will be.

Sobel Operator

Place a gradient kernel (3×3) over each pixel of the image.

Compute difference in intensity \rightarrow approximates the derivative.

Produces two output images:

- X-Direction (G_x): vertical edges
- Y-Direction (G_y): horizontal edges


Using kernel convolution, sharp changes appear as bright lines (edges).

Sobel Operator

Normalized x-gradient from Sobel operator

-1	0	+1
-2	0	+2
-1	0	+1

G_x




Sobel Operator

Normalized y-gradient from Sobel operator

+1	+2	+1
0	0	0
-1	-2	-1

G_y




Edge Detection



$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$

Applications



Applications



MIT License

Copyright (c) 2010-2017 Gert-Ludwig Ingold

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.